

Fig 7. BitMamba-2 vs BitMamba-3 at 130M (same 480M tokens, same training)

