

Fig 8. Architecture $\times$ Quantization decomposition at 130M / 480M tokens
M2 $\rightarrow$ M3 (ternary): 1.64 $\times$ advantage | FP $\rightarrow$ ternary (M3): +12.1% PPL cost

